

Fahim Shahreer

+1-647-975-9120 | fahim.shahreer@gmail.com | [fshrr](#) | in/fahim-shahreer | fahimshahreer.com

TECHNICAL SKILLS

- **ML & AI:** LangChain, LlamaIndex, PyTorch, MCP, pgvector, Arize, NeMo Guardrails, DSPy, Ollama
- **Languages & Frameworks:** Python (FastAPI), TypeScript (Next.js, React.js, Node.js), Go (net/http), SQL
- **Infrastructure & Ops:** AWS, Azure, Kubernetes, Docker, GitHub Actions, PostgreSQL, Redis

EXPERIENCE

Technical Cofounder

December 2024 — Present

Echoes Events

Toronto, CA

- Acted as fractional CTO, translating 100+ non-technical stakeholder discovery interviews into architectural requirements and full-stack production systems.
- Built a Python ELT pipeline across 400+ event sources: deduplication, normalization, LLM-based tagging, and human review to filter low-quality or off-platform events before publishing.
- Shipped a TanStack Start app on Azure covering host onboarding, event posting, and booking flows.
- Implemented embedding-based event recommendations and buddy matching to personalize discovery.

Forward Deployed AI Engineer

January 2023 — Present

Freelancer

Toronto, CA

- Deployed an air-gapped multilingual document intelligence system (LangChain, Docling, BGE-M3, LlamaIndex, Ollama) for an immigration law firm, enabling multilingual document Q&A with citations, IRCC checklist gap detection, and cross-document inconsistency flagging, cutting clerical prep time 50%.
- Shipped a voice intake pipeline (Twilio, Whisper, GPT-4o) for a physiotherapist that converts patient voicemails into structured symptom briefs delivered to a React dashboard, cutting session time from 45 to 30 minutes.
- Evaluated prompt engineered Gemini 2.5 with RAG against LoRA fine-tuned Qwen2.5-7B for accuracy, latency and serving cost when delivering a philosopher chatbot, grounded in source texts with cross-session memory and guardrails.

Full Stack Software Engineer

Jun 2022 — Jan 2023

Scotiabank

Toronto, CA

- Delivered e-transfer send and receive UI, Scotiabank's most-used digital banking feature, serving 10M+ users.
- Maintained 30+ BFF REST APIs in Next.js — trimming payloads, redacting PII, and decoupling frontend data needs.
- Created a 20+ component shared React UI library, eliminating 5% redundant code and 40% recurring bugs.

Full Stack ML Engineer

Mar 2021 — Apr 2022

Surgical Safety Technologies

Toronto, CA

- Shipped an AWS video AI pipeline processing 200+ hrs/day for 5 hospitals, with real-time inference and transcoding.
- Built 3 React/D3 dashboards (admin, user, peer) using Kinesis streams and batch data, raising video uploads 400%.
- Rolled out feature-flagged canary deploys for UIs and AI models, cutting production incidents 40%.
- Provisioned a 6-node on-prem GPU cluster for HIPAA-compliant automatic model training using docker and ansible.

Forward Deployed Engineer

Sep 2017 — Dec 2019

BlueFlower Media

Toronto, CA

- Embedded with 20+ SMB clients across diverse industries to diagnose operational bottlenecks, define technical playbooks from 0→1, and deploy production multi-tenant web products.
- Built a dockerized hosting platform on VPS, replacing managed hosting providers and cutting client hosting costs 70%.

Contract Machine Learning Engineer

May 2017 — Oct 2017

University of Toronto & LAMAS

Toronto, CA

- Fine-tuned PyTorch KittiSeg for feature segmentation in architectural facade-to-3D pipeline, saving 12hrs/project.
- Built a Django web app for ML job submission and auto load-balanced GPU workloads across AWS EC2 instances.

EDUCATION

MS in Electrical and Computer Engineering @ University of Colorado Boulder

Expected 2027

HBSc in Statistics @ University of Toronto

Class of 2021