

Fahim Shahreer

📞 +1-647-975-9120 ✉ fahim.shahreer@gmail.com 💻 fahim-shahreer 📁 fshrr 🏠 fahimshahreer.com

Technical Skills

Infrastructure: Docker, Kubernetes, Ansible, Terraform/OpenTofu, Cloudformation, Traefik, Nginx, AWS, Azure, Tailscale

Languages & APIs: Python (FastAPI, Flask), TypeScript (Node.js, React.js), Go(net/http)

ML & Data: PyTorch, Numpy, Pandas, PostgreSQL, Redis

Experience

Cofounder and Lead Engineer @ Echoes Events — Toronto, CA

December 2024 — Present

- Built and operated AWS infrastructure for a consumer events platform, owning compute, storage, and deployment configuration end-to-end.
- Implemented a Python data pipeline ingesting and normalizing 400+ event sources into a unified schema with tag classification and vector embeddings.

Freelance Platform & AI Engineer — Toronto, CA

January 2023 — Present

- Delivered production AI service integrations for clients — building and deploying RAG pipelines, tool-use APIs, and MCP servers on client infrastructure.

Full Stack Software Engineer @ Scotiabank — Toronto, CA

Jun 2022 — Jan 2023

- Maintained 30+ BFF REST APIs serving Scotiabank's 10M+ digital banking users — optimizing payloads, enforcing PII redaction, and managing frontend data contracts.

Infrastructure & ML Engineer @ Surgical Safety Technologies — Toronto, CA

Mar 2021 — Apr 2022

- Designed and operated EKS-based Kubernetes platform for a real-time surgical AI pipeline — managing EC2 node pools, Kinesis ingestion, S3 storage, and CloudFormation IaC — processing 200+ hrs/day across 5 hospitals.
- Provisioned a 6-node on-prem GPU cluster for HIPAA-compliant model training using Ansible to automate node configuration, data collection, and job scheduling.
- Built internal FastAPI services for video upload, transcoding, ML inference dispatch, and dashboard data retrieval, integrated with the EKS workload orchestration layer.
- Engineered feature-flagged canary deploy pipelines for AI models and services, reducing production incidents by 40%.

Full Stack Software Engineer @ BlueFlower Media — Toronto, CA

Sep 2017 — Dec 2019

- Built Dockerized hosting infrastructure on VPS with Nginx reverse proxy and load balancing, serving 50+ client microsites at 70% lower cost than managed hosting.

Contract Machine Learning Engineer @ University of Toronto & LAMAS — Toronto, CA

May 2017 — Oct 2017

- Designed a job scheduler for an auto-scaling AWS GPU EC2 cluster, managing workload distribution and resource allocation for distributed ML inference.

Relevant Projects

Homelab — Personal

2023 — Present

- Provisioned and manage a multi-node bare-metal Kubernetes cluster on Proxmox VE — covering node lifecycle, storage provisioning, and rolling cluster upgrades.
- Automated cluster configuration and service deployment with Ansible playbooks and Terraform modules, maintaining fully reproducible infrastructure state.
- Configured Traefik ingress with Tailscale mesh overlay for zero-trust service routing and secure remote access across the cluster.
- Deployed Prometheus/Grafana observability stack with custom dashboards for cluster health, node utilization, and workload alerting.

- Self-hosts Ollama for local LLM inference and a FOSS service stack, managing persistent volume claims and service configuration.

Education

Honours Bachelor of Science in Statistics @ University of Toronto

Class of 2021

Relevant Coursework: Systems Programming, Data Structures & Algorithms, Linear Algebra, Differential Equations

Master of Science in Electrical and Computer Engineering @ Univ. of Colorado Boulder

Expected 2027

Relevant Coursework: Linux System & Kernel Programming, Semiconductor Devices, Real-Time Embedded Systems, FPGA Design