

Fahim Shahreer

☎ +1-647-975-9120 ✉ fahim.shahreer@gmail.com 💻 fahim-shahreer 📧 fshrr 🏠 fahimshahreer.com

Technical Skills

Infrastructure & Dev Tools: AWS, Docker, Kubernetes, Ansible, Terraform, Linux, Prometheus, Grafana, PostgreSQL, Redis

Languages & Frameworks: Python, Go, TypeScript, SQL, FastAPI, Next.js/React

Machine Learning: PyTorch, Langchain, LlamaIndex, Google ADK, Numpy, Pandas

Experience

Cofounder and Lead Engineer @ Echoes Events — Toronto, CA

December 2024 — Present

- Owned a wellness events platform end-to-end, using 100+ discovery interviews to frame problems and ship features.
- Built a Python batch pipeline ingesting 400+ event sources into a unified schema with tags and embeddings.
- Shipped a TanStack Start app on AWS covering host onboarding, event posting, and booking flows.
- Implemented embedding-based event recommendations and buddy matching to personalize user discovery.

Independent Machine Learning and AI Engineer — Toronto, CA

January 2023 — Present

- Built conversational AI chatbots for clients using LangChain and Google ADK with multi-turn context-aware dialogue.
- Developed RAG pipelines over client data using LlamaIndex and pgvector for personalized knowledge retrieval.
- Created custom MCP servers to connect chatbots with client knowledge bases, enabling tool-use and live data access.

Full Stack Software Engineer @ Scotiabank — Toronto, CA

Jun 2022 — Jan 2023

- Delivered e-transfer send and receive UI, Scotiabank's most-used digital banking feature, serving 10M+ users.
- Maintained 30+ BFF REST APIs in Next.js — trimming payloads, redacting PII, and decoupling frontend data needs.
- Created a 20+ component shared React UI library, eliminating 5% redundant code and 40% recurring bugs.

Full Stack ML Engineer @ Surgical Safety Technologies — Toronto, CA

Mar 2021 — Apr 2022

- Shipped and operated AWS infrastructure for a real-time surgical video AI pipeline — EKS-based inference, EC2, Kinesis, S3, and CloudFormation IaC — processing 200+ hrs/day across 5 hospitals.
- Implemented serverless FastAPI endpoints for video upload, transcoding, ML inference, and dashboard data retrieval.
- Provisioned a 6-GPU on-prem GPU cluster for HIPAA-compliant model training, using Ansible to automate data collection.
- Engineered feature-flagged canary deploy pipelines for UIs and AI models, cutting production incidents 40%.

Full Stack Software Engineer @ BlueFlower Media — Toronto, CA

Sep 2017 — Dec 2019

- Worked with 20+ small businesses to scope needs, prototype quickly, and ship production web products end-to-end.
- Built a Dockerized hosting solution on VPS with Nginx load balancing, serving 50+ microsites at 70% lower hosting cost.

Contract Machine Learning Engineer @ University of Toronto & LAMAS — Toronto, CA

May 2017 — Oct 2017

- Fine-tuned KittiSeg in Torch for feature segmentation in architectural facade-to-3D pipeline, saving 12hrs/project.
- Designed a job scheduler for auto-scaling AWS GPU EC2 cluster handling distributed ML inference workloads.

Selected Projects

Homelab — Multi-node bare-metal Kubernetes cluster on Proxmox, using Ansible, Terraform, Tailscale, Traefik, and Prometheus/Grafana to self-host and manage Ollama and FOSS services.

Protein Modelling @ iGEM Toronto — PyRosetta protein structure modeling for CRISPR and gold binding proteins on a self-provisioned local GPU cluster.

Education

Honours Bachelor of Science in Statistics @ University of Toronto

Class of 2021

Master of Science in Electrical and Computer Engineering @ Univ. of Colorado Boulder

Expected 2028