

Fahim Shahreer

📞 +1-647-975-9120 ✉ fahim.shahreer@gmail.com 🌐 in/fahim-shahreer 📁 fshrr 🏠 fahimshahreer.com

Technical Skills

Infrastructure: Kubernetes (EKS, bare-metal), Docker, Ansible, OpenTofu, ArgoCD, Cilium, Nginx, Tailscale

Languages & APIs: Python (FastAPI, Flask), Go (net/http), TypeScript (Node.js), Bash

Observability & CI/CD: Prometheus, Grafana, Loki, Hubble, Arize AI, GitHub Actions

Storage & Data: PostgreSQL, Redis, ElasticSearch, ZFS, Longhorn

Experience

Cofounder and Lead Engineer @ Echoes Events — Toronto, CA

December 2024 — Present

- Built and operated AWS infrastructure for a consumer events platform, owning compute, storage, and deployment end-to-end.
- Implemented a Python data pipeline ingesting and normalizing 400+ event sources into a unified schema with tag classification and vector embeddings.

Freelance Platform & AI Engineer — Toronto, CA

January 2023 — Present

- Delivered production AI service integrations for clients: RAG pipelines, tool-use APIs, and MCP servers deployed on client infrastructure.
- Created custom MCP servers to connect chatbots with client knowledge bases and live data access.

Full Stack Software Engineer @ Scotiabank — Toronto, CA

Jun 2022 — Jan 2023

- Delivered e-transfer send and receive UI, Scotiabank's most-used digital banking feature, serving 10M+ users.
- Maintained 30+ BFF REST APIs in Next.js, trimming payloads, redacting PII, and decoupling frontend data needs.

Infrastructure & ML Engineer @ Surgical Safety Technologies — Toronto, CA

Mar 2021 — Apr 2022

- Designed and operated EKS-based Kubernetes platform for a real-time surgical AI pipeline, managing EC2 node pools, Kinesis ingestion, S3 storage, and CloudFormation IaC, processing 200+ hrs/day across 5 hospitals.
- Provisioned a 6-node on-prem GPU cluster for HIPAA-compliant model training using Ansible to automate node configuration, data collection, and job scheduling.
- Built serverless APIs for video upload, transcoding, ML inference dispatch, and dashboard data retrieval, integrated with the EKS workload orchestration layer.
- Engineered feature-flagged canary deploy pipelines for AI models and services, reducing production incidents by 40%.

Contract Machine Learning Engineer @ University of Toronto & LAMAS — Toronto, CA

May 2017 — Oct 2017

- Designed an auto-scaling AWS GPU EC2 job scheduler managing workload distribution for distributed ML inference.
- Built a Django app with a GPU job scheduler handling 50+ Neural Style Transfer requests/hour during ML workshops.

Relevant Projects

Homelab — Personal

2023 — Present

- Provisioned and operate a multi-node bare-metal Kubernetes cluster on Proxmox VE running Talos Linux.
- Managed all VM and LXC provisioning declaratively with OpenTofu and the bpg/proxmox provider, maintaining reproducible infrastructure state across hosts.
- Deployed ArgoCD for GitOps: K8s workloads declared in git with sync diffs and topology visualization.
- Configured Cilium as the CNI: eBPF networking, kube-proxy replacement, Gateway API routing, and Hubble observability.

Education

Master of Science in Electrical and Computer Engineering @ Univ. of Colorado Boulder

Expected 2027

Relevant Coursework: Linux System & Kernel Programming, Semiconductor Devices, Real-Time Embedded Systems, FPGA Design

Honours Bachelor of Science in Statistics @ University of Toronto

Class of 2021

Relevant Coursework: Systems Programming, Data Structures & Algorithms, Linear Algebra, Differential Equations