

Fahim Shahreer

+1-647-975-9120 | fahim.shahreer@gmail.com | linkedin.com/in/fahim-shahreer | github.com/fshrr | fahimshahreer.com

Technical Skills

Languages: JavaScript, TypeScript, Python, Go, Bash, SQL, HTML/CSS

Frontend and Backend Libraries: Node.js, Django, FastAPI, React, Next.js, TanStack, D3.js, Nginx, Jest

ML Libraries: Numpy, Pandas, PyTorch, Langchain, LlamaIndex, Ollama, Google ADK, Scikit-image, Scikit-learn

Infrastructure: Linux, AWS, GCP, Docker, Kubernetes, Ansible, PostgreSQL, DynamoDB, ElasticSearch, Redis

Tools: Claude Code, Cursor, Git, Vite

Experience

Technical Cofounder @ Echoes Events

December 2024 – Present

- Cofounded a wellness startup applying lean methodology through 3 product iterations with 100+ market research interviews, transforming Echoes from wellness edtech into a Eventbrite alternative for spiritual events and retreats.
- Built a Python microservice for LLM-assisted scraping across 400+ websites and public groups, generating keyword tags for event aggregation and recommendations with 95% accuracy.
- Developed and deployed TanStack Start full-stack platform on GCP powering event discovery, event posting, LLM chatbot, host profiles, buddy matching, recommendations, and seamless booking flows.

Freelance Technical Consultant – Toronto, CA

January 2023 – Present

- Engineered conversational chatbots for clients using LangChain, Google ADK, and foundational models (Gemini/GPT), implementing advanced prompt engineering for multi-level context-aware dialogue management.
- Built RAG pipelines ingesting internal client data via LlamaIndex embeddings into PostgreSQL (pgvector), developing custom MCP servers to connect chatbots with these personalized knowledge bases.

Full Stack Software Developer @ Scotiabank – Toronto, CA

Jun 2022 – Jan 2023

- Developed and maintained over 30 Backend for Frontend REST APIs in Next.js for e-transfer functionalities, serving over 10 million users, as part of Scotiabank's Web Modernization project.
- Proposed and led the development of a shared React.js UI library for the monorepo incorporating over 20 components, leading to a 5% code reduction in the repo and a 40% decrease in recurring bugs in multiple teams.

Full Stack ML Software Developer @ Surgical Safety Technologies – Toronto, CA

Mar 2021 – Apr 2022

- Architected key features on a distributed video processing and AI analysis pipeline on AWS that handled over 200 hours of daily video uploads, transcoding, and AI analysis for 5 major US and Canadian hospitals.
- Increased the average user task completion rate by 400% through the development of 3 crucial dashboard UIs (admin, user and peer assessment dashboard) using React.js and D3.js
- Reduced user-reported bugs by 40% and cut system downtime by 100%, by proposing and integrating feature flags for canary testing the deployment of new features.

Full Stack Software Developer @ BlueFlower Media – Toronto, CA

Sep 2017 – Dec 2019

- Bootstrapped a digital media startup, leading a cross-functional team of 4 developers and 2 designers to deliver comprehensive digital solutions for over 20 small businesses, using Node.js, React.js and Wordpress.
- Designed and deployed a containerized load balanced web hosting infrastructure for low-traffic sites using Kubernetes and Nginx reverse proxy, reducing hosting costs by up to 70%.

Contract Machine Learning Engineer @ University of Toronto & LAMAS – Toronto, CA

May 2017 – Oct 2017

- Built an ML pipeline to transform 2D building models into 3D using a retrained convolutional neural network in PyTorch and Scikit-image for precise vectorization, reducing manual labeling by 12 hours per project.
- Developed a Django application with a job scheduler that load balanced AWS GPU EC2 instances to manage over 50 Neural Style Transfer job requests per hour in workshops.

Projects

test project @ group name (github.com)

May 2017 – Oct 2017

- Built an ML pipeline to transform 2D building models into 3D using a retrained convolutional neural network in PyTorch and Scikit-image for precise vectorization, reducing manual labeling by 12 hours per project.
- Developed a Django application with a job scheduler that load balanced AWS GPU EC2 instances to manage over 50 Neural Style Transfer job requests per hour in workshops.

Test project

May 2017 – Oct 2017

- Built an ML pipeline to transform 2D building models into 3D using a retrained convolutional neural network in PyTorch and Scikit-image for precise vectorization, reducing manual labeling by 12 hours per project.
- Developed a Django application with a job scheduler that load balanced AWS GPU EC2 instances to manage over 50 Neural Style Transfer job requests per hour in workshops.

Test project @ group name

May 2017 – Oct 2017

- Built an ML pipeline to transform 2D building models into 3D using a retrained convolutional neural network in PyTorch and Scikit-image for precise vectorization, reducing manual labeling by 12 hours per project.
- Developed a Django application with a job scheduler that load balanced AWS GPU EC2 instances to manage over 50 Neural Style Transfer job requests per hour in workshops.

Education

Honours Bachelor of Science in Statistics @ University of Toronto – (GPA: 9001/9000)

Class of 2021

Relevant Courses: Data Science, Computer Science, Machine Learning

Honours Bachelor of Science in Statistics @ University of Toronto

Class of 2021

Relevant Courses: Data Science, Computer Science, Machine Learning